

Stories built with care, not just code.



An updated public guide to the product, demonstrated foundation, open Sprint 1 gate, Child Map Discovery, Story Studio architecture, roadmap, design language, and useful unknowns.

VERIFIED + CANDID

16 PAGES

Same design DNA. More mature architecture. No false finish line.

The project moved forward without rewriting the past.

Sprint 0 is demonstrated. Sprint 1 has substantial implementation and evidence, but its outcome gate remains explicitly open. Sprint 2 and Sprint 3 may proceed only inside a documented synthetic-input exception.

OUTCOME DEMONSTRATED

Sprint 0

Clean React and .NET foundations, architecture, CI, contracts, dependency checks, and browser evidence passed the recorded outcome gate.

GATE OPEN

Sprint 1

Family access and protected-parent-mode work is implemented deeply. A required provider proof remains unresolved, so the gate is not passed.

A temporary sequencing exception allows synthetic Sprint 2/3 work. It does not rename Sprint 1 as complete, and the unresolved evidence must be revisited before Sprint 4.

A parent turns care and intent into a story a child can enjoy.

Tiny Custom Stories aims to make personalized learning playful and understandable while keeping the technology bounded and worthy of family trust.

ROLE 1

Parents lead

Choose purpose, context, review, and approval. The system should not silently decide what family information belongs in a story.

ROLE 2

Children receive wonder

The first audience is ages two through six, reading with a parent or independently from a parent-approved library.

ROLE 3

AI assists

Generation can help plan and write. It does not own family authority, Story truth, or the decision to make a story child-visible.

The product is not trying to maximize content or attention. It is trying to help a family make a meaningful story for a purpose: morals and values, fun and fantasy, or reading development.

A completed sprint is a demonstrated outcome.

Sprint 0 passed only after the same identified foundation was checked locally and in CI, with unresolved risks carried forward rather than hidden.

01	Buildable foundations	React/Vite and ASP.NET Core foundations build and test independently of the retired applications.
02	Contracts	The checked-in API contract reproduces deterministically and agrees with runtime behavior.
03	Quality gates	Formatting, linting, types, analyzers, Release builds, tests, and dependency thresholds passed.
04	Browser evidence	A controlled browser journey and reviewed artifacts showed the web and API foundations working together.
05	Clean hand-off	The recorded gate found no remaining Sprint 0 must-have work.

A strong implementation can still have an open outcome gate.

Sprint 1 built the family-access boundary deeply: child-safe default mode, protected parent authority, expiry, recovery, isolation, protected work restoration, accessibility, and end-to-end evidence. One required provider proof remains unresolved.

BOUNDARY 1

Child mode

Default after sign-in. Read-only. Parent data and protected actions stay unavailable.

BOUNDARY 2

Parent mode

Deliberate adult entry with server-enforced authority for sensitive family and story work.

BOUNDARY 3

Awkward states

Expiry, refresh, lock contention, recovery, and concealed protected work are part of the behavior.

BOUNDARY 4

Gate status

Substantial implementation is not being used as a substitute for the unresolved provider evidence.

The awkward moments were treated as product behavior.

The gate remains open, but the evidence base is much richer than the September 2 public snapshot showed.

01	Family isolation	First and repeat sign-in preserve one internal family boundary without cross-family access.
02	Child-safe authorization	Child mode cannot read protected parent information or successfully call parent-only APIs.
03	Expiry and recovery	Entry, failure, expiry, refresh, lock contention, and recovery paths have targeted coverage.
04	Protected work	Temporary parent-only work restores only inside the authorized boundary.
05	Accessibility	Keyboard, focus, announcements, reflow, error recovery, and responsive evidence are explicit.
06	End-to-end journey	Synthetic browser evidence exercises sign-in, protected entry, and return to child mode.

The Child Map became a living record, not a completion form.

Sprint 2 separates durable Child Map information from Discovery operating state. Discovery can evolve without hard-coding one questionnaire into the main application.

DISCOVERY 1

Versioned questions

Approved question definitions are content with stable identifiers and answer schemas.

DISCOVERY 2

Typed answers

The renderer understands the question shape rather than baking each prompt into a bespoke screen.

DISCOVERY 3

Deterministic eligibility

Age, context, history, pacing, and cooldown rules decide what may be asked before optional future intelligence.

DISCOVERY 4

Bounded follow-ups

Follow-up depth is explicit and limited instead of allowing an open-ended hidden conversation.

Sprint 2 does not send Child Map text or photos to an external AI provider for runtime interpretation. Future AI authoring, ranking, or phrasing can sit behind replaceable capability boundaries.


Story creation became a reversible composition workflow.

The canonical Sprint 3 experience is one guided Story Studio from intent through approval. The story is a composition, not a one-shot prompt response.

01	Review context	The parent sees a small permitted personalization set, may remove items, add story-only details, or choose a completely fictional mode.
02	Plan before writing	Story purpose and reading choices become a visible Story Plan before generation.
03	Compose ten text pages	Sprint 3 proves a cover plus exactly ten text pages. AI-generated illustrations remain Sprint 4.
04	Revise locally	Direct editing, paragraph-scoped assistance, page rewrite, visible continuity choices, and lightweight restore reduce destructive regeneration.
05	Approve a stable version	Approval creates Story truth that later draft edits do not silently mutate.

Separate responsibility before separating machines.

The project now makes ownership explicit inside the application first. A capability may later become a worker or service when operational evidence justifies that boundary.



Story capability ownership

Story Lifecycle owns Story truth, persistence, revisions, approval, and deletion participation.

Story Context Selection owns policy, permitted candidates, parent selection, and ContextBundle construction.

Story Generation consumes minimized versioned artifacts and returns candidates/findings; it does not own family authorization or Story persistence.

Physical extraction remains evidence-led. Worker/service extraction, durable workflow technology, broker, deployment platform, and production MCP adapters are still open operational choices.

Seven outcomes, with the current exception visible.

The roadmap remains outcome-based. The temporary Sprint 1 sequencing exception permits bounded synthetic Sprint 2/3 work; it does not convert the open Sprint 1 gate into a pass.

00	Foundation	DEMONSTRATED
01	Accounts + parent mode	GATE OPEN
02	Child Map + Discovery	CANONICAL
03	Story Studio	CANONICAL
04	Characters + art	LATER
05	Library + reader	LATER
06	Alpha readiness	LATER

Before Sprint 4 starts, the unresolved Sprint 1 provider evidence must be revisited. Later sprint definitions remain proposals until their scope and evidence mature.

Detailed input is part of the engineering system.

The project intentionally prefers precise, reviewable implementation slices over a small number of ambiguous tickets. This also makes AI-assisted implementation less dependent on rediscovering architecture and product intent.

STEP 1

Human direction

Product and architecture choices are made explicitly and recorded.

STEP 2

Explicit tasks

Dependencies, non-goals, acceptance criteria, tests, and visual references are written down.

STEP 3

AI-assisted implementation

Implementation agents work inside the documented contracts rather than inventing hidden product decisions.

STEP 4

Evidence + PR

Automated checks, targeted boundary evidence, review, and the sprint outcome gate determine whether work is done.

Task count is not a success metric. Clarity, reviewability, evidence, and safe dependency order are more important than keeping the backlog artificially small.

The why behind every screen, service, and hand-off.

The same philosophy now spans child, parent, public, documentation, and technical work without forcing every surface into one visual template.

Human intention leads

AI assists inside parent-led purpose and review.

Two modes, two needs

Child and parent authority and density remain distinct.

Calm over capture

No attention traps, autoplay, streaks, or artificial urgency.

Warmth with clarity

Character supports meaning instead of competing with it.

Trust is visible

Privacy, uncertainty, drafts, progress, and limits appear where they matter.

Every state is designed

Loading, interruption, error, expiry, recovery, and return count.

Accessibility is structural


Semantics, focus, contrast, reflow, alternatives, and reduced motion begin early.

Evidence outranks taste

Synthetic inspected evidence decides whether confidence is deserved.

Same design DNA, more mature story.

The public site keeps its paper-and-ink editorial identity. New architecture and development content extends existing components and rhythms rather than introducing a generic dashboard or documentation-site aesthetic.



INK
#27213B



PAPER
#FBF8EF



PLUM
#6F549F



CORAL
#E76042



SUN
#F7B84B



MINT
#9DD3C6



LAVENDER
#E7D8F5

Display: Fraunces. Body and controls: Nunito Sans. Metadata: DM Mono. The generated dossier keeps accepted fallback fonts when those web fonts are unavailable.

Continuity rule: existing foundation -> additive evolution -> targeted refinement. New content should feel like the September site grew, not like somebody replaced it.

Useful unknowns stay visible.

The capability boundaries are clearer, but privacy, derived-artifact lifecycle, provider transfer, physical extraction, and Alpha evidence still require deliberate decisions.

Q01

Which Child Map information is necessary, optional, or prohibited?

Q02

Which permitted Child Map observations may be used for a specific story or external provider?

Q03

What happens to drafts and approved stories when their Child Map sources are later deleted?

Q04

When should Story Generation move from the API process to a worker or service?

Q05

Which providers, workflow products, and deployment choices meet the evidence bar?

Q06

What proves the private Alpha is valuable and safe enough to continue?

Certainty deserves a label. Privacy deserves a boundary.

The public story is a curated translation. It does not render internal documentation automatically or expose private project links, family information, credentials, or attack-relevant operating detail.

CONFIRMED DECISION

A deliberately adopted choice.

VERIFIED FACT

Supported by named repository evidence.

PROPOSAL

A candidate direction, not a promise.

ASSUMPTION

A belief that still needs evidence.

OPEN QUESTION

An unresolved matter kept visible.

Every material public update should preserve knowledge status, public safety, existing route compatibility, and the site's established visual language unless newer accepted evidence explicitly supersedes it.

Keep the boundaries clear enough to change safely.

The near-term project is not one giant autonomous build. It is a chain of explicit product decisions, small implementation slices, public-safe architecture, synthetic evidence, review, and outcome gates.

FOCUSED NEXT WORK

Finish the evidence

Revisit the unresolved Sprint 1 provider proof before Sprint 4.

FOCUSED NEXT WORK

Build Discovery safely

Implement accepted versioned question and Child Map lifecycle contracts with synthetic data first.

FOCUSED NEXT WORK

Compose stories reversibly

Build Story Studio around policy-limited context, scoped revisions, and stable approval.

FOCUSED NEXT WORK

Extract only when earned

Move generation into a worker/service only when operational evidence justifies the added boundary.

Source trail: Sprint 0 outcome evidence; Sprint 1 specification and evidence; DEC-2026-039 through DEC-2026-043; Sprint 2 and Sprint 3 canonical specifications; ADR-0016 and ADR-0017; current architecture; design philosophy; visual design language; and public-story curation guidance.